



Campy AI

PLAIN ENGLISH HANDBOOK

How your cameras start paying attention

What we connect to, what the system watches
for, and what you will actually see — explained
without a single piece of jargon.

Plain English Handbook

01	The problem this solves	Your cameras already record everything. Nobody watches.
02	What we connect to	Your existing cameras. No new hardware.
03	How the camera is connected	What happens on installation day
04	How it actually works	From a picture to a message, in six steps
05	The seven things it watches for	And what each one is genuinely good at
06	What you will see	Alerts, the daily view, and who gets told what
07	Why it does not spam you	The part most systems get wrong
08	Teaching it about your site	It gets better because you correct it
09	Privacy and your staff	What is stored, what is not, and what to tell people
10	Honest limits	What it cannot do
11	Common questions	The ones we are always asked

The problem this solves

Your cameras already record everything that happens. The trouble is that footage only helps *after* something has gone wrong — and only if somebody is willing to scrub through hours of it.

A control room with forty screens has one person watching. Within twenty minutes, nobody is really watching any of them. This is not a failure of attention; it is what happens to every human being asked to stare at mostly nothing.

So the fire is found when somebody smells it. The restricted door being propped open is found when there is an accident. The stock going missing is found at the next count, three weeks later.

Campy AI watches those same camera feeds continuously and tells you the handful of things that actually matter, as they happen, in ordinary sentences:

"Unauthorised entry into Machine Guard Zone — Assembly Line A — 21:14."

Not a number on a chart. Not a graph you need an analyst to read. A sentence you can act on.

What changes for you

TODAY	WITH CAMPY AI
You find out after the event	You find out while it is happening
Somebody has to be watching	Nobody has to be watching
Reviewing footage takes hours	The relevant ninety seconds is already picked out
Rules depend on supervision	The same rule applies at 3 a.m. and at shift change

What we connect to

Your existing cameras. That is the whole answer. You do not buy new cameras, you do not re-cable anything, and nothing about your current setup is thrown away.

Will it work with what I have?

Almost certainly yes. If your cameras were installed in the last ten years or so, they can already send video over your network, and that is all we need. This includes:

- Any modern IP camera — Hikvision, Dahua, CP Plus, Axis, Uniview, and the dozens of brands that resell them.
- Cameras plugged into an NVR or DVR box, which is the usual arrangement. We connect to the box.
- Older analogue cameras, provided they run into a recorder. The recorder does the converting; the cameras themselves never change.

A simple test

If you can already see your cameras on a phone app or a web page, Campy AI can read them too. It is the same video, arriving the same way.

What we do need

WHAT	WHY
The camera or recorder's address on your network	So the software knows where to look. Your installer or IT person has this.
The username and password for the cameras	The same ones used to view them today.
One computer on the same network	This is where Campy AI runs. It can be a server in a cupboard; it does not need to be anything exotic.

What we do not need

- **No new cameras.** The point of the product is that you already own the hardware.
- **No cloud upload.** Your video never leaves your building. There is nothing to upload and no external service watching it.
- **No internet connection for the cameras.** In fact we recommend the opposite — cameras should not be reachable from outside at all.
- **No change to how you record.** Your existing recording keeps working exactly as it does now. Campy AI watches alongside it.

How the camera is connected

Installation day, start to finish. Most sites are working within an afternoon.

01 The software is installed on one computer

A server on the same network as the cameras. Nothing is installed on the cameras themselves, and nothing is installed on staff computers — people use it through a web browser.

02 Your sites and cameras are listed

A "site" is a building or a floor. Each camera is added under its site with a name a person would use — *Main Entrance, Loading Bay, Stockroom* — not a serial number.

03 Each camera is pointed at

Usually this is just the camera's address and password. The system then asks the camera where its video is, and the camera answers. There is a **Test connection** button that either shows a picture or tells you plainly what is wrong.

04 You choose what each camera watches for

Not every camera needs everything. The stockroom camera watches for theft and objects left behind. The corridor camera watches crowding. The camera over the machinery watches the guard zone. Switching off what a camera does not need makes it faster and quieter.

05 Areas are drawn on the picture

Somebody drags a shape around the machinery on a still image from that camera and calls it *Machine Guard Zone*. That is a geofence — a line the system understands. It takes about fifteen seconds per area.

06 You say who gets told what

Fire goes to everybody, immediately. Crowding at the entrance goes to the floor supervisor during working hours. A restricted-area entry at 2 a.m. goes to the security manager's phone. Each rule is a sentence you fill in.

You can see it working before you connect anything

Campy AI includes built-in demonstration cameras that play realistic simulated scenes — someone entering a restricted area, a fire starting, a corridor filling up. You can watch every feature work and decide whether you want it before an installer visits your site.

How the video travels

The Campy AI computer reaches out to each camera and pulls pictures from it, several times a second. Nothing is sent the other way, and no camera is reachable from the internet. Your staff see the results in a browser, the same way they would use any internal website.

How it actually works

Six steps, several times a second, for every camera. Nothing here is magic, and it helps to know what the system is really doing.

01 It takes a picture

About six still pictures every second from each camera. That is plenty — nobody walks into a restricted room and out again in a sixth of a second.

02 It learns what the empty room looks like

Over the first minute or so it builds up a picture of the room with nobody in it: the shelves, the floor, the machinery. From then on it can tell what has changed. It keeps updating, so when somebody moves a pallet, the pallet becomes part of the room within a minute rather than being reported forever.

03 It works out what the changed bits are

Is that shape a person? A box? A trolley? This is where the actual AI happens — a model that has been trained on what people and objects look like from a camera's point of view.

04 It follows each person from picture to picture

This is the difference between "there is a person" and "*this* person walked from the door to the shelf and stood there for ninety seconds." Without it you have counting; with it you have behaviour.

05 It checks each rule you switched on

Is anyone inside a restricted area? How many people are in this corridor? Does that flickering orange patch look like fire? Is there a bag that has sat still for four minutes with nobody near it?

06 It decides whether a human needs to know

Most of what it notices is not worth interrupting anybody for. What survives becomes an alert, with the picture attached, in a sentence that says what happened, where, and when.

The honest version of "AI"

It is not watching in the way a person watches. It has no idea what your business does, and it never decides that somebody is guilty of anything. What it does is notice specific patterns very reliably, at 3 a.m., on forty cameras at once, without ever getting bored. That is a genuinely useful thing to have, and it is worth understanding it as exactly that rather than as something smarter.

The seven things it watches for

Each one can be switched on or off per camera. Here is what each is genuinely good at — and where each is weaker.

Movement and posture

Whether people are walking, standing, sitting, bending, running, or have not moved in a long time. It flags somebody running indoors, or somebody who appears to have fallen and stayed down.

Good for: spotting a fall in an area nobody passes through often; noticing unsafe speed near machinery. **Weaker at:** telling you *why* — it knows somebody ran, not whether they were late or fleeing.

Recognising your staff

If you choose to enrol them, the system can say *who* rather than *somebody*. This is off until you switch it on, and requires written consent per person.

Good for: knowing whether a restricted-area entry was an authorised engineer or a stranger. **Weaker at:** poor light, faces turned away, or people in masks.

Restricted areas

You draw a line around an area; the system tells you when somebody crosses it who should not. Areas can be restricted only at certain hours.

Good for: server rooms, machinery, stockrooms, roof access — the single most useful feature on most sites. **Weaker at:** nothing much. This is the most reliable of the seven.

Crowding

How many people are in an area, and whether the flow is orderly or turning into a crush.

Good for: entrances at shift change, canteens at lunch, campus corridors between classes. **Weaker at:** very dense crowds where people are mostly

hidden behind each other, where it estimates rather than counts.

Theft-related behaviour

Repeated handling of stock, unusual carrying, lingering at shelves, items disappearing from a shelf area.

This flags behaviour, never people

It says *this is worth a look* — never *this person is a thief*. It cannot see inside a bag or a pocket, and it does not know what your stock is worth. Treat every one of these as a prompt to check the footage yourself, and say so in your staff policy.

Dangerous or abandoned objects

Things that should not be where they are: a tool left in a walkway, a bag nobody is standing near, a spill nobody has cleared.

Good for: unattended items in public areas; obstructions in fire escapes.

Weaker at: naming the object. It is often more certain that *something was left* than about *what* it is — which is usually the part that matters anyway.

Fire and smoke

Flame colour, the specific way fire flickers, and a fire's tendency to grow — checked together, then confirmed by a trained model.

Good for: high ceilings, loading bays, open yards and warehouse racking, where a ceiling-mounted smoke sensor is slow or useless. **Weaker at:** fires outside the camera's view, and smoke behind racking. It complements your fire alarm; it does not replace it, and it is not a certified fire detection system.

What you will see

Three things, depending on who you are and what you came to do.

An alert, when something happens

An email, a message, or a notice in the app — whichever you chose. It contains a sentence, the camera name, the time, and the picture from the moment it happened, with the relevant part outlined.

The daily view

When somebody opens Campy AI, the first screen answers four questions without being asked: what happened in the last day, what still needs somebody's attention, whether every camera is working, and how quickly the team is responding to alerts.

The live wall

Every camera at once, with what the system currently sees drawn on top — the people it is following, the areas you have marked, the counts. This is the screen for a control room.

Different people see different things

Everybody signs in as themselves, and what they can see and do follows their job:

ROLE	WHAT THEY GET
Owner	Everything, including the bill.
Administrator	Runs the system day to day.
Security manager	Cameras, restricted areas, incidents, who gets alerted.
Control-room operator	The live wall and incoming alerts. Nothing else — deliberately.
HR manager	Staff-related insight. No camera configuration.
Auditor	Read-only, plus the full record of who did what.
Viewer	Look, do not touch.

You can also build your own roles if none of these match how your site is organised.

Why it does not spam you

This is the part most camera-analytics systems get wrong, and the reason most of them end up switched off within a month.

A system that sends four hundred alerts a day is worse than no system at all. People stop reading them, and then the one that mattered goes unread too. So a detection has to survive three separate filters before anybody is disturbed:

01 The same thing seen repeatedly is one thing

Somebody standing in a restricted area appears in every picture for as long as they stand there. That is one event with a counter, not sixty events.

02 Related events become one incident

A person crossing three zones on their way through a building is one incident to review, not three separate things to work out the connection between.

03 Your rules decide who is actually told

You set how serious something has to be, how sure the system has to be, how long to stay quiet after alerting once, how many times something must happen before it counts at all, and which hours the rule applies. You can also say: if nobody acknowledges this within ten minutes, tell somebody more senior.

And when it is wrong, you say so

Every alert has a **false alarm** button. That is not just closing it — that picture can be fed back into the system as an example of what *not* to flag, and the mistake is trained out. The system gets quieter over time because your team corrects it, not because somebody turned down a sensitivity dial and lost real detections along with the noise.

Teaching it about your site

A system trained on pictures from the internet has never seen your uniform, your machinery, or your idea of unsafe. Campy AI can be taught.

The AI in Campy AI is entirely ours — we built the whole thing rather than renting somebody else's. The practical consequence for you is that you can train it yourself, on your own pictures, on your own computer, and the result belongs to you.

What that looks like in practice

01 Collect examples

Upload pictures, take them from a camera, or — best of all — use the alerts your team marked as false alarms. Those are already labelled by the people who know the site.

02 Label them in the browser

Click through pictures saying what each one is. A few hundred is usually enough. No software to install, no data scientist required.

03 Press train, and watch

A chart shows the accuracy climbing while it learns. A small model takes seconds to a few minutes.

04 Put it to work, or undo it

Deploy the new version to your cameras with one click. If it turns out worse than the old one, go back to the previous version with one click. Every version is kept with its score.

What this is genuinely useful for

- Teaching it your high-visibility vests, so it can tell staff from visitors.
- Teaching it the specific thing that keeps causing false fire alerts on one camera — a reflection, a heater, a particular light.
- Teaching it equipment specific to your industry that no general model has ever seen.

Privacy and your staff

Monitoring people is a responsibility, not a feature. How you introduce this matters as much as how it is configured.

What the system does and does not store

WHAT HAPPENS	
Video	Not stored by Campy AI at all. Your existing recorder keeps doing that. Campy AI keeps still pictures of moments that triggered an alert.
Faces	If you enrol staff, what is stored is a list of numbers describing the face — not a photograph. You cannot rebuild a face from it. Deleting it is one click.
Everyone else	Faces of people not enrolled can be automatically blurred in every stored picture. This is a setting you can switch on for the whole system.
Old records	Deleted automatically on a schedule you set — thirty days for pictures by default.
Who looked	Every sensitive action is written to a record that cannot be quietly edited, and it exports.

Face recognition stays off until you turn it on

It requires written consent recorded against each individual. Without that consent, the system will not enrol them. This is deliberate friction.

What to tell your staff

Our strong advice, from sites where this has gone well and sites where it has not:

- **Tell them before you switch it on**, not after somebody notices.
- **Be specific about what it watches for**. "Restricted areas, crowding and fire" reassures. "AI monitoring" alarms.
- **Say plainly that it is not a productivity monitor**. It does not score people, rank them, or report how long they spent away from a workstation.
- **Explain that theft detection flags behaviour for review**, and that a human always decides what it means.

- **Put the retention period in writing.** People are far more comfortable when they know pictures are deleted after thirty days.

The strongest argument for your staff

On most sites the honest pitch is safety, not surveillance. It notices a fall in a room nobody walks past. It notices a fire while it is small. It notices a crush forming at an exit. Those are things that protect the people being watched.

Honest limits

Every system of this kind has limits. A vendor who does not tell you theirs is hiding them.

It can only see what the camera sees

A camera pointed at the wrong place, with a dirty lens, or facing into the afternoon sun, gives poor results. Camera placement matters more than any setting in the software.

People need to be reasonably large in the picture

Somebody at the far end of a large warehouse may be too small to detect reliably. This is a matter of physics, not software.

It does not understand intent

It sees that someone entered a restricted area. It does not know whether they are an engineer with permission or an intruder — unless you have enrolled your staff, in which case it can tell you who.

Theft detection does not detect theft

It detects behaviour that is often associated with it. A person is always the one who decides what actually happened.

Fire detection is not a fire alarm

It is early warning that complements your alarm system. It is not a certified replacement and must not be treated as one.

Very dark scenes are hard

If your cameras do not have usable night vision, the system sees roughly what you would see on the same footage — which is not much.

Watching costs computing power

Every camera being analysed takes real processing. Plan for two to three cameras per processor core with everything switched on — more if a camera only needs one or two of the seven. Your installer measures this on the actual machine before committing to a camera count.

Common questions

Do we have to buy new cameras?

No. That is the entire point of the product. If your cameras can be viewed on a phone or a computer today, they will work.

Does our video go to the internet?

No. Everything runs on a computer in your building. There is no third-party AI service in the path and nothing to upload. We recommend your cameras have no internet access at all.

Will this slow down or interfere with our recording?

No. Campy AI reads the same stream your recorder reads; it does not take it over. Your existing recording continues untouched. Very old cameras that allow only two viewers at a time are the one thing to check with your installer.

How long does installation take?

Most sites are running the same day. The software installs in minutes; the time goes into adding cameras, drawing areas and agreeing who should be alerted about what.

Can we try it before connecting real cameras?

Yes. Built-in demonstration cameras play realistic simulated scenes, so you can watch every feature work before an installer visits.

The camera says the password is wrong, but it is right

Your camera's address is often given as one long line that already contains the username and password. If you paste that *and* type them into the login boxes as well, the camera is sent the login twice and refuses it. Put the plain address in the address box and the login in the login boxes — Campy AI now separates them for you if you paste the long version.

Can I get a report I can print or send on?

Yes — for one feature or for all of them together, over whatever period you choose. It arrives as a PDF containing the picture behind every finding, with the person or object outlined and the area they crossed drawn on it, so a safety meeting or an insurer can see the moment itself rather than take your word for it. Download it from the Analytics page.

What happens if a camera goes offline?

The system notices, retries with increasing patience, and marks the camera degraded and then offline. Camera health is on the front screen, so a camera that has quietly stopped working does not go unnoticed for weeks.

How many cameras can one server handle?

With all seven features on, plan for two to three cameras per processor core. Cameras that only need one or two features, or that watch quiet areas, cost much less. Your installer measures your actual hardware rather than guessing.

Will it flood us with alerts?

It is specifically designed not to. Three separate filters stand between a detection and a message, and every alert has a false-alarm button that teaches the system. See chapter 7.

Can it tell us who someone is?

Only if you enrol them and record their written consent. It is switched off by default.

Is this legal where we operate?

Workplace monitoring rules vary by country and often by state. The system gives you what compliance needs — consent gating, automatic deletion, face blurring and a full audit trail — but the decision to monitor, and the notice you give staff, is yours. Take local advice.

What does it cost to run?

A subscription, plus the electricity for one computer. There is no per-camera cloud fee and nothing metered per picture analysed — which is precisely why the AI is ours rather than rented.

What if we want to stop?

Uninstall the software. Your cameras, your recorder and your footage are exactly as they were — Campy AI never changed any of them.
